



Instruction Level Parallelism (ILP)

(Micro-Architecture)

A Computer System



MIPS R4000 Pipeline



Topics in Computer Architecture

- Instruction set
- Registers, Memory Map, Addressing modes
- Binary arithmetic (2's complement, IEEE 754 floating point standard, addition, multiplication, division), from DLD
- Input and output
 - Memory Hierarchy (Virtual memory, Main Memory, Cache)
 - Datapaths (single-cycle, multicycle, pipeline)
 - Control (combinational logic, FSM, microcode)
 - Pipelining (throughput, hazards, forwarding, stall, branch prediction)
 - Advanced architectures (ILP, OOE, superscalar, etc.)
 - Multiprocessors, Parallel Processing
 - GPGPUs
 - Performance (benchmarks, energy efficiency, Amdal's law)
 - Not discussed in this course:
 - Power management

Advanced Architectural Concepts

- Can we achieve $CPI < 1$?
(i.e., can we have $IPC > 1$?)

State-of-the-Art Microprocessor

- “Superscalar” execution or Instruction Level Parallelism (ILP)
+ “Deeper Pipeline $\Rightarrow$ Dynamic Branch Prediction $\Rightarrow$ Speculation $\Rightarrow$ Recovery
- “Out-of-order” (OOE) Execution $\Rightarrow$ Instruction Buffer and Prefetch $\Rightarrow$ Reorder Buffers
- “VLIW” Very Long Instruction Word Ex: Intel/HP Titanium


Instruction Parallelism (ILP)

- **Key idea:** Allow instructions behind stall to proceed
- **Instruction level parallelism (ILP):** multiple instructions fetched and executed simultaneously
- Enables **out-of-order** execution $\Rightarrow$ out-of-order completion
 - ID stage checks for hazards. If no hazards, issue the instn for execution. (**Scoreboard ILP** dates to CDC 6600 in 1963)

Key Ideas– ILP

- ILP is used **in addition to pipelining**.
- Processors with ILP are called **multiple-issue processors** — **multiple instructions launched in 1 clock cycle**. Two ways:
 - MIMD: Multiple Instructions Multiple Data
 - **Scalar Processors** {
 - Very long instruction word (VLIW) — **static** multiple issue
 - Super-pipeline
 - Superscalar — **dynamic** multiple issue
 - SIMD: Single Instruction Multiple Data
 - **Vector processor**


Super-pipeline and Superscalar



How ILP Works

- Issuing multiple instructions per cycle would require fetching multiple instructions from memory per cycle => called **Superscalar degree** or **Issue width**
- To find independent instructions, we must have a big pool of instructions to choose from, called **instruction buffer (IB)**. As IB length increases, complexity of decoder (control) increases that increases the datapath cycle time
- **Prefetch** instructions sequentially by an Instruction Fetch Unit (IFU) that operates independently from datapath control. Fetch instruction **(PC)+L**, where L is the IB size or as directed by the branch predictor.

Microarchitecture of a Superscalar CPU (Power PC)



Microarchitecture of a Superscalar CPU (Intel Pentium 4)




Figure 4: Pentium® 4 processor microarchitecture

A Static Two-Issue MIPS Pipeline

- ILP determined by **Compiler**
- Read **two instructions per cycle**
 - **First** : An ALU or branch instruction, **and**
 - **Second** : A load or store instruction
- Insert one **NOP** if above pair is not possible for Exec
- **Added hardware**
 - A **second instruction memory (Buffer/Queue)**
 - Additional **input/output ports** in register file
 - **Additional ALU** in execute stage for address calculation

An Example

```
Loop:    lw      $t0, 0($s1)
         addu    $t0, $t0, $s2
         sw      $t0, 0($s1)
         addi    $s1, $s1, -4
         bne     $s1, $0, Loop
```

With pipeline
CPI = 1
Total = 5

Instruction Queue/Buffer

Static Two-Issue Execution

	ALU or branch instruction	Data transfer instruction	Clock cycle
Loop:	nop	lw \$t0, 0(\$s1)	1
	addi \$s1, \$s1, -4	nop	2
	addu \$t0, \$t0, \$s2	nop	3
	bne \$s1, \$0, Loop	sw \$t0, 4(\$s1)	4

Note *code reordering to avoid pipeline stalls and change in sw argument (4) since addi now before sw nop since data dependency*

$$\text{CPI} = \frac{4}{5} = 0.8 > 0.5 \text{ (ideal without nop)}$$



 With pipeline

With Loop Unrolling (Multiple of 4)

Loop Unrolling is essential for ILP Processors **Why?**

But, increase in Code memory and no. of registers.

	ALU or branch instruction	Data transfer instruction	Clock cycle
Loop:	addi \$s1, \$s1, -16	lw \$t0, 0(\$s1)	1
	nop	lw \$t1, 12(\$s1)	2
Note Register Change By Compiler	addu \$t0, \$t0, \$s2	lw \$t2, 8(\$s1)	3
	addu \$t1, \$t1, \$s2	lw \$t3, 4(\$s1)	4
	addu \$t2, \$t2, \$s2	sw \$t0, 16(\$s1)	5
	addu \$t3, \$t3, \$s2	sw \$t1, 12(\$s1)	6
	nop	sw \$t2, 8(\$s1)	7
	bne \$s1, \$0, Loop	sw \$t3, 4(\$s1)	8

$$\text{CPI} = 8/14 = 0.57 > 0.5 \text{ (ideal)}$$

With pipeline

VLIW: Very Long Instruction Word


- Static multiple issue, ILP determined by compiler.
- Datapath contains multiple execution units.
- Compiler groups instructions that have no data or resource conflicts for parallel execution.
- Grouped instructions are packed in very long words of a wide instruction memory.
- Speedup benefit of VLIW is highly program dependent.

J. A. Fisher, “Very Long Instruction Word Architecture and ELI-512,” *Proc. 10th Symp. on Computer Architecture*, Stockholm, June 1983, pp. 478-490.

J. A. Fisher, P. Faraboschi and C. Young, *Embedded Computing: A VLIW Approach to Architecture, Compilers and Tools*, Morgan Kaufmann.

Very Large Instruction Word (VLIW)

$IPC > 1$, $CPI < 1$




VLIW Architectures

- Wide-issue processor that **relies on compiler** to
 - Packet together independent instructions to be issued in parallel
 - Schedule code to minimize hazards and stalls
- Very long instruction words (3 to 8 operations)
 - Can be issued in parallel without checks
 - If compiler cannot find independent operations, it inserts **nops**
- Advantage: simpler HW for multiple issue (ILP)
 - Faster clock cycle
- Disadvantages:
 - Code size
 - Requires aggressive compilation technology

VLIW and Superscalar


- Sequential stream of long instruction words
- Instructions scheduled statically by the compiler
- Number of simultaneously issued instructions is fixed during compile-time
- Instruction issue is less complicated than in a **superscalar processor**
- **Disadvantage:** VLIW processors cannot react on dynamic events, e.g. cache misses, with the same flexibility like superscalars.
- Padding VLIW instructions with no-ops is needed in case the full issue bandwidth is not be met. This increases code size. More recent VLIW architectures use a denser code format which allows to remove the no-ops.
- *VLIW is an architectural technique, whereas superscalar is a microarchitecture technique.*

Traditional VLIW Hardware



- Multiple functional units, many registers (e.g. 128)
 - Large multiported register file (for N-FUs need $\sim 3N$ ports)
- Simple instruction fetch unit
 - No checks, direct correspondence between slots & FUs
- Instruction format
 - 16 to 24 bits per op $\Rightarrow 5 \times 16 = 80$ bits to $5 \times 24 = 120$ bits wide

TriMedia TM32 Architecture



Compressed code in the Instruction Cache

VLIW Code Example

for (i=0; i<100; i++) X[i] = X[i] + B;

- MIPS code

```
Loop:      l.d f2, 0(r1)
           add.d f4, f0, f2
           s.d f4, 0(r1)
           addui r1, r1, 8
           bne r1, r2, Loop
```

Unrolled Loop (7 times) by Compiler VLIW

#	Mem1	Mem2	FP1	FP2	Int/Branch
1	l.d f2, 0(r1)	l.d f6,8(r1)			
2	l.d f10,16(r1)	l.d f14,24(r1)			
3	l.d f18,32(r1)	l.d f22,40(r1)	add.d f4,f0, f0	add.d f8,f0, f6	
4	l.d f26,48(r1)		add.d f12,f0,f10	add.d f16,f0,f14	addu r1,r1,#56
5			add.d f20,f0,f18	add.d f24,f0,f22	
6	s.d f4,-56(r1)	s.d f8,-48(r1)	add.d f28,f0,f26		
7	s.d f12,-40(r1)	s.d f16,-32(r1)			
8	s.d f20,-24(r1)	s.d f24,-16(r1)			
9	s.d f28, -8(r1)				bneq r1,r2,Loop

- Assuming 2 cycle loads, 3 cycle FP adds (with pipeline)
- 9 cycles for 7 iterations
 - 1.28 cycles/iteration
 - 2.6 operations/cycle
- 21 nops (empty slots)

23 24 FP registers used (instead of just 3)

Problems with Early VLIW

- Characteristics
 - Large number of slots per long instruction
 - All operations in a long instructions must be independent
 - No interlocks (check for hazards across long instructions)
 - Fixed/direct correspondence between instruction slots and FUs
- Code size problems
 - Need many nops within each long instructions
 - Potential solution: **code compression**
- Binary compatibility problems
 - Different ISA, every time number or latency of FUs changes
 - **Must Recompile for every processor generation**
 - Potential solution: **binary translation**

Modern VLIW Processors

- Hardware interlocks
 - HW check for dependencies across LIWs
- Use a special field (bit mask) to determine if operations in LIW are independent or not
 - E.g. stop bits in IA-64
 - Can also specify dependence to following LIW
- Support several combinations of operation types per LIW
 - E.g., allow both Mem-FP-Int and Mem-Int-Int
 - Requires network that will “route” operations to proper FUs


How Can the HW Help the Compiler with Discovering ILP?

- Compiler's performance is critical for VLIW processors
 - Find many independent instructions & schedule them in best possible way
- What limits the compiler's ability to discover ILP
 - Name dependencies (WAW & WAR)
 - Can eliminate with large number of registers
 - Branches
 - Limit compiler's ability to schedule
 - Modern VLIW processors use branch prediction too
 - Dependencies through memory
 - Force the compiler to use conservative schedule
- Can the HW help the compiler?
 - Superscalar processors

Superscalar Design of P4 (CISC)

- CISC shell:
 - Processor **fetches** instructions from memory **in the order** of static program.
 - Each instruction is **translated** into one or more fixed-length RISC instructions, known as micro-operations (micro-ops).
- RISC core (**Execution**):
 - Micro-ops are **executed out-of-order** in a dynamically scheduled pipeline.
 - Processor **commits the result** of each micro-op execution to register file **in order** of original program flow.

Superscalar: Dynamic Scheduling and Out-of-Order Execution



Superscalars

- 3 or more instruction issues per clock:
 - Intel P6
 - AMD K5
 - Sun UltraSPARC
 - Alpha 21164
 - MIPS R10000
 - PowerPC 604/620
 - HP 8000
- References:
 - D. W. Anderson, F. J. Sparacio and R. M. Tomasulo, “The IBM 360 Model 91: Processor Philosophy and Instruction Handling,” *IBM J. Res. Dev.*, vol. 11, pp. 8-24, January 1967.
 - T. Agerwala and J. Cocke, “Reduced Instruction Set Processors,” Technical Report RC12434 (#55845), Yorktown Heights, NY: IBM T. J. Watson Research Center, January 1987.

Out of Order Execution (OOE)

- A procedural programming language sequences instructions.
- Sequencing assumes an order of execution – no parallelism.
- **OOE** must preserve correctness of result.
- **Principle**: Two instructions can be executed in parallel if they do not have dependencies.

RAW Dependence

- Read after write (RAW): A dependent instruction reads from a register being written to by another instruction.
- Example:

add **\$s1**, \$s2, \$s3

sub \$s2, **\$s1**, \$s3

sub has RAW dependence on **add**

WAR Dependence

- Write after read (WAR): A dependent instruction writes to a register being read by another instruction.
- Example:

add \$s1, **\$s2**, \$s3

sub **\$s2**, \$s0, \$s3

sub has WAR dependence on **add**

WAW Dependence

- Read after write (RAW): One instruction writes to a register to being written to by another instruction.
- Example:

add **\$s2**, \$s2, \$s3

.....

sub **\$s2**, \$s1, \$s3

sub has WAW dependence on **add**

Superscalar Instruction Issue

- Rules:

- RAW dependence — If any operand is being written, **do not issue**.
- WAR dependence — If the result register is being read, **do not issue**.
- WAW dependence — If the result register is being written, **do not issue**.

- Scoreboard:

- Cycle by cycle record of registers and execution units showing how many instructions are using them.
- Example 1: In-order issue (next 2 slides).
- Example 2: Out-of-order issue (3rd slide).

Dynamic Scheduling (Scoreboard)

- Consider an example:
 - First with in-order issue
 - Then with out-of-order issue
- Assume:
 - Up to two instructions are **fetch**ed in a cycle
 - Instruction register can hold two instructions
 - An Instruction is **issu**ed in decode cycle, **or must wait** until there is no RAW, WAR or WAW dependence
 - An instruction can **retire two or three cycles** after it is issued

Ck cycle	Inst #	Decoded	Issue Inst#	Retire Inst#	Reg. to read								Reg. to write							
					0	1	2	3	4	5	6	7	0	1	2	3	4	5	6	7
1	1	R3 = R0 * R1	1		1	1										1				
	2	R4 = R0 + R2	2		2	1	1									1	1			
2	3	R5 = R0 + R1	3		3	2	1									1	1	1		
	4	R6 = R1 + R4	-		3	2	1									1	1	1		
3					3	2	1									1	1	1		
4		In Order Issue		1	2	1	1										1	1		
				2	1	1												1		
				3																
5		*****	4			1			1										1	
	5	R7 = R1 * R2	5			2	1		1										1	1
6	6	R1 = R0 - R2	-			2	1		1										1	1
7				4		1	1													1
8				5																
9		*****	6		1		1													
	7	R3 = R3 * R1	-		1		1							1						

In-order Issue scoreboard (Continued)

Ck cycle	Instr #	Decoded	Issue Inst#	Retire Inst#	Reg. to read								Reg. to write							
					0	1	2	3	4	5	6	7	0	1	2	3	4	5	6	7
10					1		1							1						
11				6																
12	8	R1 = R4 + R4	7			1		1								1				
			-			1		1								1				
13						1		1								1				
14						1		1								1				
15				7																
16			8						2					1						
17									2					1						
18				8																

Out-of-order scoreboard (Next 2 Slides)

Questions?

- RAW dependence: Inst# 4 ($R6 = R1 + R4$) could not be issued until cycle 5. Should Inst# 5 ($R7 = R1 * R2$) wait in queue?
- Answer: No. Inst# 5 can be issued in cycle 3 as there is no register conflict (out-of-order issue).
- WAR dependence: Must the issue of Inst#6 ($R1 = R0 - R2$) wait until cycle 9 when all instructions reading R1 have retired?
- Answer: No. Provided new result of Inst#6 does not affect R1 being used by previous instructions (register renaming: $R1 \rightarrow R1^*$).

Ck cycle	Inst #	Decoded	Issue Inst#	Retire Inst#	Reg. to read								Reg. to write																												
					0	1	2	3	4	5	6	7	0	1	2	3	4	5	6	7																					
1	1	R3 = R0 * R1	1		1	1											1																								
	2	R4 = R0 + R2	2		2	1	1										1	1																							
2	3	R5 = R0 + R1	3		3	2	1										1	1	1																						
	4	R6 = R1 + R4	-		3	2	1										1	1	1																						
3	5	R7 = R1 * R2	5	2	3	3	2										1	1	1																						
	6	S1 = R0 - R2	6		4	3	3										1	1	1																						
					3	3	2									1		1																							
4	7 8	R3 = R3 * S1 S2 = R4 + R4	4 - 8	1 3	3	4	2		1								1		1	1																					
					3	4	2		1							1		1	1																						
					3	4	2		3							1		1	1																						
					2	3	2		3								1	1	1																						
					1	2	2		3									1	1																						
5				6		2	1		3							1			1																						
6			7	4 5 8		2	1	1	3					1		1			1	1																					
																					1	1	1	1	1	1	1	1	1	1	1										
																																1	1	1	1	1	1	1	1	1	1
7									1								1																								
8									1								1																								
9				7																																					

μ-architecture: Superscalar

Two types: Scoreboard and Tomasulo

Scoreboard (EX: PENTIUM):

- **Out-of-order execution** divides **ID stage**:
 1. **Issue**—decode instructions, check for **structural hazards**
 2. **Read operands**—wait until no **data hazards**, then read operands
- **Scoreboards** allow instruction to execute whenever there is no structural hazard or not waiting for prior instructions. So the instructions are **issued in order**, but can bypass the waiting instructions in the read operand stage => **In-order issue Out-of-Order execution => Out-of-Order completion**
- Named after **CDC 6600 Scoreboard**, which developed this capability

Scoreboard Implications

- Scoreboard replaces ID, EX, WB with 4 stages
- Out-of-order completion => WAR, WAW hazards?
- Solutions for WAR => Wait at the WB stage until the other instruction completes
- For WAW, must detect hazard at the ID stage: stall until other completes
- Need to have multiple instructions in execution phase => multiple execution units or pipelined execution units
- Scoreboard keeps track of dependencies, state or operations

CDC 6600 Scoreboard

- Speedup 1.7 from compiler; 2.5 by hand
BUT slow memory (no cache) limits benefit
- Limitations of 6600 scoreboard:
 - No forwarding hardware
 - Limited to instructions in basic block (small *window*)
 - Small number of functional units (structural hazards), especially integer/load store units
 - Do not issue on structural hazards
 - Wait for WAR hazards
 - Prevent WAW hazards

Summary of ILP + Scoreboard

- Instruction Level Parallelism (ILP) in SW or HW
- Loop level parallelism is easiest to see
- SW parallelism dependencies defined for program, hazards if HW cannot resolve
- SW dependencies/compiler sophistication determine if compiler can unroll loops
 - Memory dependencies hardest to determine
- HW exploiting ILP
 - Works when can't know dependence at run time
 - Code for one machine runs well on another
- Key idea of Scoreboard: Allow instructions behind stall to proceed (Decode $\Rightarrow$ Issue instr & read operands)
 - Enables out-of-order execution $\Rightarrow$ out-of-order completion
 - ID stage checked both for structural

Tomasulo Algorithm

(Implemented in IBM 360/91 in 1966)

- Control & buffers distributed with Function Units (FU) vs. centralized in scoreboard;
 - FU buffers called “reservation stations”; have pending operands
- Registers in instructions replaced by values or pointers to reservation stations(RS); called register renaming ;
 - avoids WAR, WAW hazards
 - More reservation stations than registers, so can do optimizations compilers can't
- Results to FU from RS, not through registers, over Common Data Bus that broadcasts results to all FUs
- **Load and Stores** treated as FUs with RSs as well

Tomasulo Summary

- **Reservations stations**: renaming to larger set of registers + buffering source operands
 - Prevents registers as bottleneck
 - Avoids WAR, WAW hazards of Scoreboard
 - Allows loop unrolling in HW
- Helps cache misses as well
- Lasting Contributions
 - Dynamic scheduling
 - **Register renaming**
 - Load/store disambiguation
- **IBM 360/91 descendants** are Pentium II; PowerPC 604; **MIPS R10000**; HP-PA 8000; Alpha 21264

Tomasulo v. Scoreboard

IBM 360/91 v. CDC 6600

Pipelined Functional Units

(6 load, 3 store, 3 +, 2 x/÷)

window size: ≤ 14 instructions

No issue on structural hazard

WAR: renaming avoids

WAW: renaming avoids

Broadcast results from FU

distributed reservation stations

Multiple Functional Units

(1 load/store, 1 +, 2 x, 1 ÷)

≤ 5 instructions

same

stall completion


stall completion

Write/read registers

central scoreboard

HW Schemes: Out-of-order execution


- Out-of-order execution divides ID stage:
 1. **Issue**—decode instructions, check for structural hazards, **Issue in order** if the functional unit is free and no WAW.
 2. **Read operands (RO)**—wait until no data hazards, then read operands
 - ADDB would stall at RO, and SUBB could proceed with no stalls.



- Scoreboards allow instruction to execute whenever 1 & 2 hold, **not waiting for prior instructions.**

HW support: Reorder Buffer


- Need HW buffer for results of uncommitted instructions: *reorder buffer*
 - 3 fields: instr, destination, value
 - Reorder buffer can be operand source => more registers like RS
 - Use reorder buffer number instead of reservation station when execution completes
 - Supplies operands between execution complete & commit
 - Once operand commits, result is put into register
 - Instructions commit in order
 - As a result, its easy to undo speculated instructions on mispredicted branches or on exceptions



HW support for More ILP

- **Speculation**: allow an instruction to issue that is dependent on branch predicted to be taken *without any consequences* (including exceptions) if branch is not actually taken (“HW undo”); called “boosting” *branch predicted ++
- Combines branch prediction with dynamic scheduling to execute before branches resolved
- Separate *speculative* *bypassing of results* from real bypassing of results
 - When instruction no longer speculative, write boosted results (instruction commit) or discard boosted results
 - execute out-of-order but commit in-order to prevent irrevocable action (update state or exception) until instruction commits


Branch Instruction Speculation



Renaming Registers

- Common variation of *speculative* design
- Reorder buffer keeps instruction information but not the result
- Extend register file with extra renaming registers to hold speculative results
- Rename register allocated at issue;
result into rename register on execution complete;
rename register into real register on commit
- Operands read either from register file
(real or speculative) or via Common Data Bus
- *Advantage*: operands are always from single source
(extended register file)

Speculative, Out-of-Order Superscalar Processor



Power Reduction by Slack Scheduling

- Application: Superscalar, out-of-order execution:
 - An instruction is executed as soon as the required data and resources become available.
 - A commit unit reorders the results.
- Delay the completion of instructions whose result is not immediately needed.
- Example of RISC instructions:
 - add r0, r1, r2; (A)
 - sub r3, r4, r5; (B)
 - and r9, r1, r9; (C)
 - or r5, r9, r10; (D)
 - xor r2, r5, r11; (E)

J. Casmira and D. Grunwald,
“Dynamic Instruction Scheduling
Slack,” *Proc. ACM Kool Chips
Workshop*, Dec. 2000.

Slack Scheduling

